

# Résumé du TP2 - Programmation avec l'API HDFS

## I. Démarrer le Cluster Hadoop

### 1. Démarrage des containers

```
user@Sara-N:~$ docker start hadoop-master hadoop-slave1 hadoop-slave2
hadoop-master
hadoop-slave1
hadoop-slave2
```

### 2. Accéder au master

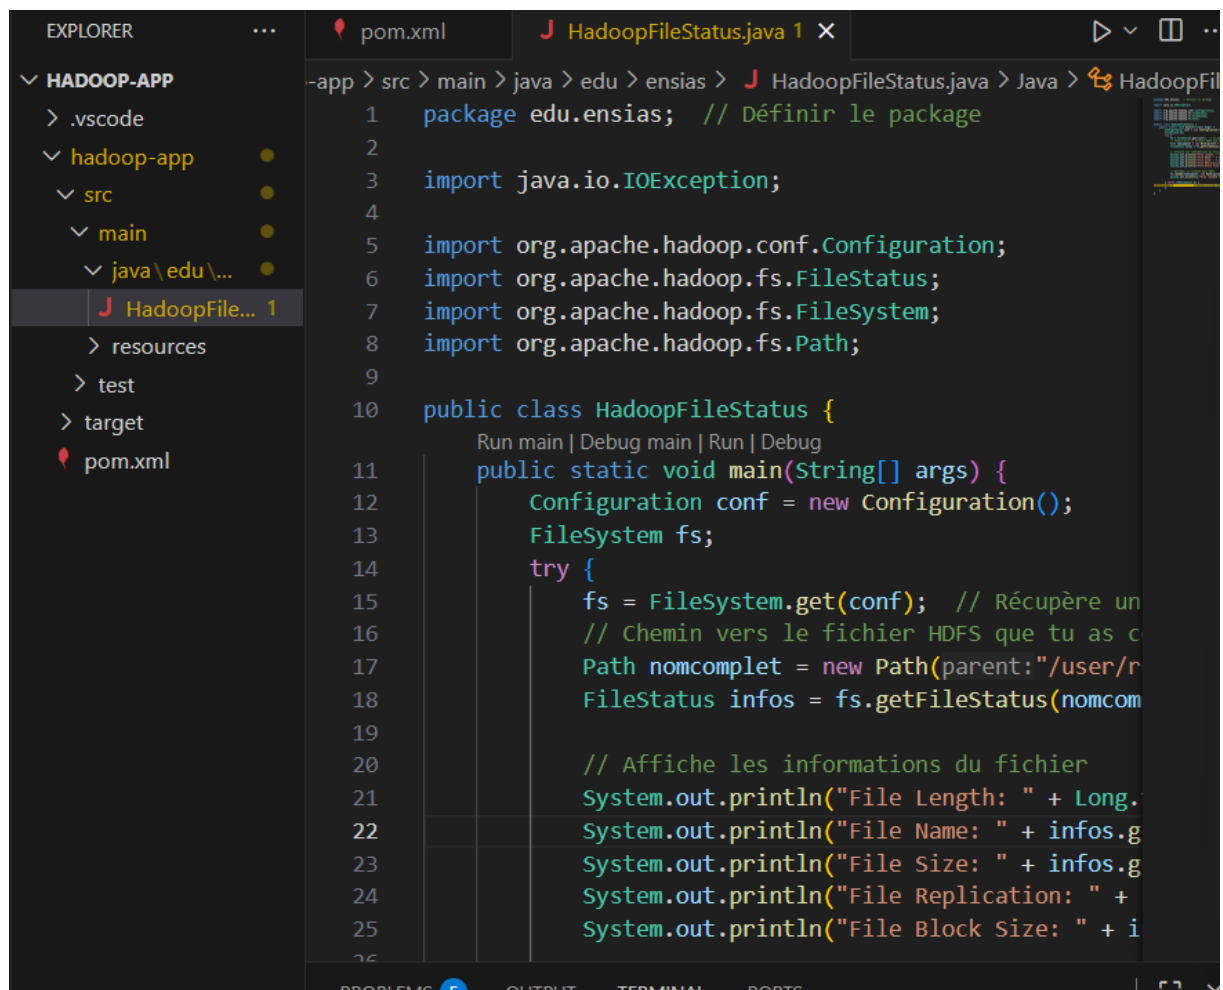
```
user@Sara-N:~$ docker exec -it hadoop-master bash
root@hadoop-master:~#
```

## II. Programmation avec l'api HDFS

### 1. Installation de l'environnement de développement

```
root@hadoop-master:~# java -version
openjdk version "1.8.0_362"
OpenJDK Runtime Environment (build 1.8.0_362-8u372-ga~us1-0ubuntu1~18.04-b09)
OpenJDK 64-Bit Server VM (build 25.362-b09, mixed mode)
```

### 2. Premier exemple



- Créer un fichier jar qu'on va nommer HadoopFileStatus.jar

```
PS C:\Users\USER\Documents\workspace_vscode\hadoop-app\hadoop-app> mvn clean package
```

[INFO] Building jar: C:\Users\USER\Documents\workspace\_vscode\hadoop-app\hadoop-app\target\hadoop-app-1.0-SNAPSHOT.jar

Puis j'ai renommé manuellement le jar 'HadoopFileStatus.jar' .

- Copier le jar créé vers le dossier de partage /hadoop\_project

```
PS C:\Users\USER\Documents\workspace_vscode\hadoop-app\hadoop-app> docker cp target\HadoopFileStatus.jar
hadoop-master:/shared_volume/
Successfully copied 5.12kB to hadoop-master:/shared_volume/
```

```
root@hadoop-master:~# ls -l /shared_volume
total 160
-rwxr-xr-x 1 root root 3169 Oct 12 16:19 HadoopFileStatus.jar
-rw-r--r-- 1 root root 2549 Oct 7 00:06 achat.txt
-rwxr-xr-x 1 1000 1000 150886 Oct 7 14:42 alice.txt
-rwxr-xr-x 1 1000 1000 2549 Oct 6 23:36 purchases.txt
```

- Lancer la commande

```
root@hadoop-master:~# hadoop jar /shared_volume/HadoopFileStatus.jar
File Length: 2549 octets
File Name: purchases.txt
File Size: 2549
File Replication: 2
File Block Size: 134217728
File renamed to achats.txt
```

- Modifier la classe HadoopFileStatus pour qu'elle puisse lire les paramètres chemin\_fichier nom\_fichier nouveau\_nom\_fichier lors de l'exécution

```
public class HadoopFileStatus {
    Run main | Debug main | Run | Debug
    public static void main(String[] args) {
        // Vérification des arguments
        if (args.length != 3) {
            System.out.println("Usage: HadoopFileStatus <chemin_dossier_HDFS> <nom
            System.exit(status:1);}

        String cheminDossier = args[0]; // ./input
        String nomFichier = args[1];    // purchases.txt
        String nouveauNom = args[2];    // achats.txt

        Configuration conf = new Configuration();
        try {
            FileSystem fs = FileSystem.get(conf);

            // Chemin complet du fichier source
            Path cheminCompleet = new Path(cheminDossier, nomFichier);

            // Récupération des informations du fichier
            FileStatus infos = fs.getFileStatus(cheminCompleet);
            System.out.println("File Length: " + infos.getLength() + " octets");
            System.out.println("File Name: " + infos.getPath().getName());
            System.out.println("File Size: " + infos.getLength());
            System.out.println("File Replication: " + infos.getReplication());
            System.out.println("File Block Size: " + infos.getBlockSize());
```

Building jar: C:\Users\USER\Documents\workspace\_vscode\hadoop-app\hadoop-app\target\HadoopFileStatus.jar

```
root@hadoop-master:~# hdfs dfs -ls /user/root/input
Found 1 items
-rw-r--r--  2 root supergroup      2549 2025-10-12 17:02 /user/root/input/purchases.txt
```

Exécuter le jar:

```
root@hadoop-master:~# hadoop jar /shared_volume/HadoopFileStatus.jar /user/root/input purchases.txt achats.txt
File Length: 2549 octets
File Name: purchases.txt
File Size: 2549
File Replication: 2
File Block Size: 134217728
File renamed to achats.txt
```

Input après:

```
root@hadoop-master:~# hdfs dfs -ls /user/root/input
Found 1 items
-rw-r--r--  2 root supergroup      2549 2025-10-12 17:02 /user/root/input/achats.txt
```

### 3. Info d'un fichier sur HDFS

```
public class HDFSInfo {
    Run main | Debug main | Run | Debug
    public static void main(String[] args) throws IOException {
        if (args.length != 1) {
            System.out.println(x:"Usage: HDFSInfo <chemin_fichier_HDFS>");
            System.exit(status:1); }

        String fichierHDFS = args[0]; // ex: /user/root/input/achats.txt

        Configuration conf = new Configuration();
        FileSystem fs = FileSystem.get(conf);

        Path cheminCompleet = new Path(fichierHDFS);

        if (!fs.exists(cheminCompleet)) {
            System.out.println(x:"Le fichier n'existe pas");
        } else {
            FileStatus infos = fs.getFileStatus(cheminCompleet);
            System.out.println("File Length: " + infos.getLength() + " octets");
            System.out.println("File Name: " + infos.getPath().getName());
            System.out.println("File Replication: " + infos.getReplication());
            System.out.println("File Block Size: " + infos.getBlockSize());

            // Afficher les blocs
            BlockLocation[] blocks = fs.getFileBlockLocations(infos, start:0, infos.getLength());
            System.out.println("Number of blocks: " + blocks.length);
            for (int i = 0; i < blocks.length; i++) {
                System.out.println("Block " + i + ": " + blocks[i].getOffset() + " -> " + (blocks[i].getOffset() + blocks[i].getLength()));
            }
        }

        fs.close();
    }
}
```

```
PS C:\Users\USER\Documents\workspace_vscode\hadoop-app\hadoop-app> mvn clean package
```

Building jar: C:\Users\USER\Documents\workspace\_vscode\hadoop-app\hadoop-app\target\HDFSInfo.jar

```
docker cp target/HDFSInfo.jar hadoop-master:/shared_volume/
>>
Successfully copied 6.66kB to hadoop-master:/shared_volume/
```

```
root@hadoop-master:~# ls -l /shared_volume
total 168
-rwxr-xr-x 1 root root  4747 Oct 12 17:10 HDFSInfo.jar
-rwxr-xr-x 1 root root  3261 Oct 12 16:45 HadoopFileStatus.jar
-rw-r--r-- 1 root root  2549 Oct  7 00:06 achat.txt
-rwxr-xr-x 1 1000 1000 150886 Oct  7 14:42 alice.txt
-rwxr-xr-x 1 1000 1000  2549 Oct  6 23:36 purchases.txt
```

```

root@hadoop-master:~# hadoop jar /shared_volume/HDFSInfo.jar /user/root/input/achats.txt
File Length: 2549 octets
File Name: achats.txt
File Replication: 2
File Block Size: 134217728
Number of blocks: 1
Block 0: 0 -> 2549

```

#### 4. Lire un fichier sur HDFS

```

public class HDFSInfo {
    Run main | Debug main | Run | Debug
    public static void main(String[] args) throws IOException {
        if (args.length != 1) {
            System.out.println("Usage: HDFSInfo <chemin_fichier_HDFS>");
            System.exit(status:1); }

        String fichierHDFS = args[0]; // ex: /user/root/input/achats.txt

        Configuration conf = new Configuration();
        FileSystem fs = FileSystem.get(conf);
        Path cheminCompleto = new Path(fichierHDFS);
        if (!fs.exists(cheminCompleto)) {
            System.out.println("Le fichier n'existe pas");
        } else {
            FileStatus infos = fs.getFileStatus(cheminCompleto);
            System.out.println("File Length: " + infos.getLength() + " octets");
            System.out.println("File Name: " + infos.getPath().getName());
            System.out.println("File Replication: " + infos.getReplication());
            System.out.println("File Block Size: " + infos.getBlockSize());
            // Afficher les blocs
            BlockLocation[] blocks = fs.getFileBlockLocations(infos, start:0, infos.getLength());
            System.out.println("Number of blocks: " + blocks.length);
            for (int i = 0; i < blocks.length; i++) {
                System.out.println("Block " + i + ": " + blocks[i].getOffset() + " -> " + (blocks[i].getOffset() + blocks[i].getLength()));
            }
        }
        fs.close();
    }
}

```

Building jar: C:\Users\USER\Documents\workspace\_vscode\hadoop-app\hadoop-app\target\ReadHDFS.jar

```

PS C:\Users\USER\Documents\workspace_vscode\hadoop-app\hadoop-app> docker cp target/ReadHDFS.jar hadoop-master:/shared_volume/
Successfully copied 7.68kB to hadoop-master:/shared_volume/

```

```

root@hadoop-master:~# ls -l /shared_volume
total 176
-rwxr-xr-x 1 root root 4747 Oct 12 17:10 HDFSInfo.jar
-rwxr-xr-x 1 root root 3261 Oct 12 16:45 HadoopFileStatus.jar
-rwxr-xr-x 1 root root 5891 Oct 12 17:29 ReadHDFS.jar
-rw-r--r-- 1 root root 2549 Oct 7 00:06 achat.txt
-rwxr-xr-x 1 1000 1000 150886 Oct 7 14:42 alice.txt
-rwxr-xr-x 1 1000 1000 2549 Oct 6 23:36 purchases.txt

```

```

root@hadoop-master:~# hadoop jar /shared_volume/ReadHDFS.jar /user/root/input/achats.txt
2012-01-01 09:00 San Jose Men's Clothing 214.05 Amex
2012-01-01 09:00 Fort Worth Women's Clothing 153.57 Visa
2012-01-01 09:00 San Diego Music 66.08 Cash
2012-01-01 09:00 Pittsburgh Pet Supplies 493.51 Discover
2012-01-01 09:00 Omaha Children's Clothing 235.63 MasterCard
2012-01-01 09:00 Stockton Men's Clothing 247.18 MasterCard
2012-01-01 09:00 Austin Cameras 379.6 Visa
2012-01-01 09:00 New York Consumer Electronics 296.8 Cash
2012-01-01 09:00 Corpus Christi Toys 25.38 Discover
2012-01-01 09:00 Fort Worth Toys 213.88 Visa
2012-01-01 09:00 Las Vegas Video Games 53.26 Visa
2012-01-01 09:00 Newark Video Games 39.75 Cash
2012-01-01 09:00 Austin Cameras 469.63 MasterCard
2012-01-01 09:00 Greensboro DVDs 290.82 MasterCard
2012-01-01 09:00 San Francisco Music 260.65 Discover
2012-01-01 09:00 Lincoln Garden 136.9 Visa
2012-01-01 09:00 Buffalo Women's Clothing 483.82 Visa

```

## 5. Ecrire un fichier sur HDFS

```
public class WriteHDFS {  
    Run main | Debug main | Run | Debug  
    public static void main(String[] args) throws IOException {  
        if (args.length != 2) {  
            System.out.println("Usage: WriteHDFS <chemin_fichier_HDFS> <text>");  
            System.exit(status:1);  
        }  
  
        String fichierHDFS = args[0]; // ex: /user/root/input/bonjour.txt  
        String contenu = args[1];     // texte à écrire dans le fichier  
  
        Configuration conf = new Configuration();  
        FileSystem fs = FileSystem.get(conf);  
        Path cheminCompleto = new Path(fichierHDFS);  
  
        if (!fs.exists(cheminCompleto)) {  
            FSDataOutputStream outStream = fs.create(cheminCompleto);  
            outStream.writeUTF(contenu); // écrit le texte dans le fichier  
            outStream.close();  
            System.out.println("Fichier créé et texte écrit sur HDFS : " + fichierHDFS);  
        } else {  
            System.out.println("Le fichier existe déjà : " + fichierHDFS);  
        }  
        fs.close();  
    }  
}
```

Building jar: C:\Users\USER\Documents\workspace\_vscode\hadoop-app\hadoop-app\target\WriteHDFS.jar

```
PS C:\Users\USER\Documents\workspace_vscode\hadoop-app\hadoop-app> docker cp target/WriteHDFS.jar hadoop-master:/shared_volume/  
Successfully copied 8.7kB to hadoop-master:/shared volume/
```

```
root@hadoop-master:~# hadoop jar /shared_volume/WriteHDFS.jar /user/root/input/bonjour.txt "Bonjour tout le monde !"  
Fichier cr?? et texte ?crit sur HDFS : /user/root/input/bonjour.txt
```

```
root@hadoop-master:~# hdfs dfs -ls /user/root/input  
Found 2 items  
-rw-r--r--  2 root supergroup      2549 2025-10-12 17:02 /user/root/input/achats.txt  
-rw-r--r--  2 root supergroup       25 2025-10-12 17:39 /user/root/input/bonjour.txt
```